Machine Learning

Squad Rapidminer

Francisco Ávila

Rubén Tenreiro

Marco Hernani

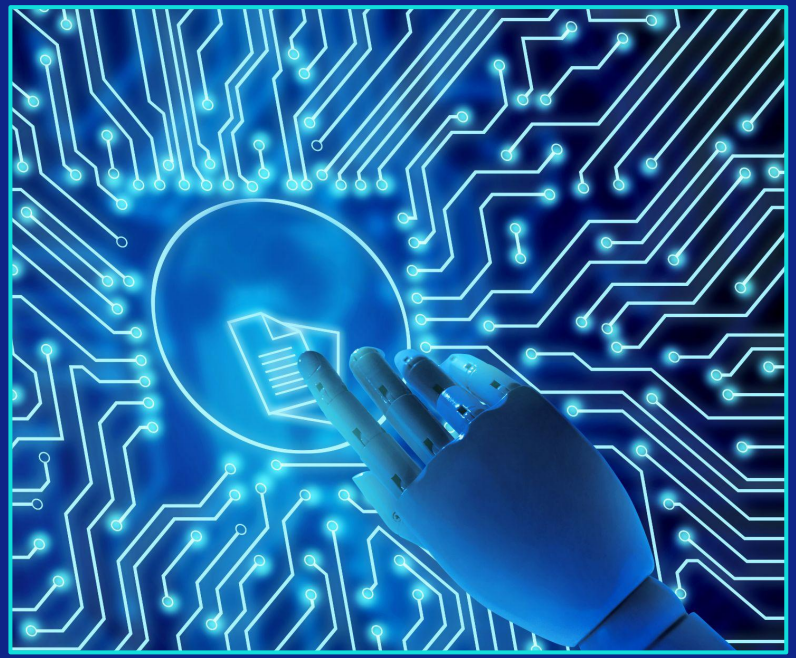

INDICE

https://archive.ics.uci.edu/ml/datasets/Census+Income

01

Introducción

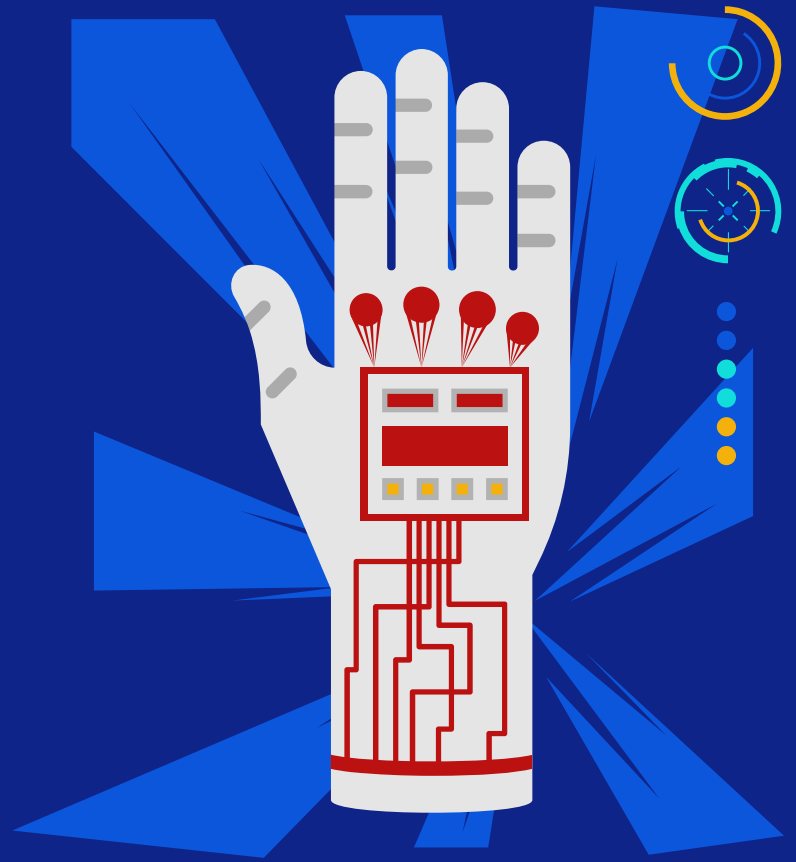

Variables cuantitativas

age	Edad
fnlwgt	Número que el censo cree que la entrada representa
education_num	Un número relacionado con el nivel de educación (sin un orden específico)
capital_gain	Capital ganado con inversiones
capital_loss	Capital perdido en inversiones
hours_per_week	Horas trabajadas por semanas

Variables cualitativas

work_class	Clase en la que se categoriza el trabajo de la persona
education	Nivel más alto de educación de la persona
marital_status	Estado civil de la persona
occupation	Trabajo de la persona
relationship	Tipo de relación que tiene esa persona
race	Raza de la persona
native_country	País de nacimiento de la persona
sex	Sexo de la persona

Variable objetivo

target	Variable objetivo que contiene un valor binario en función de si una persona gana o no más de 50K

02

EDA

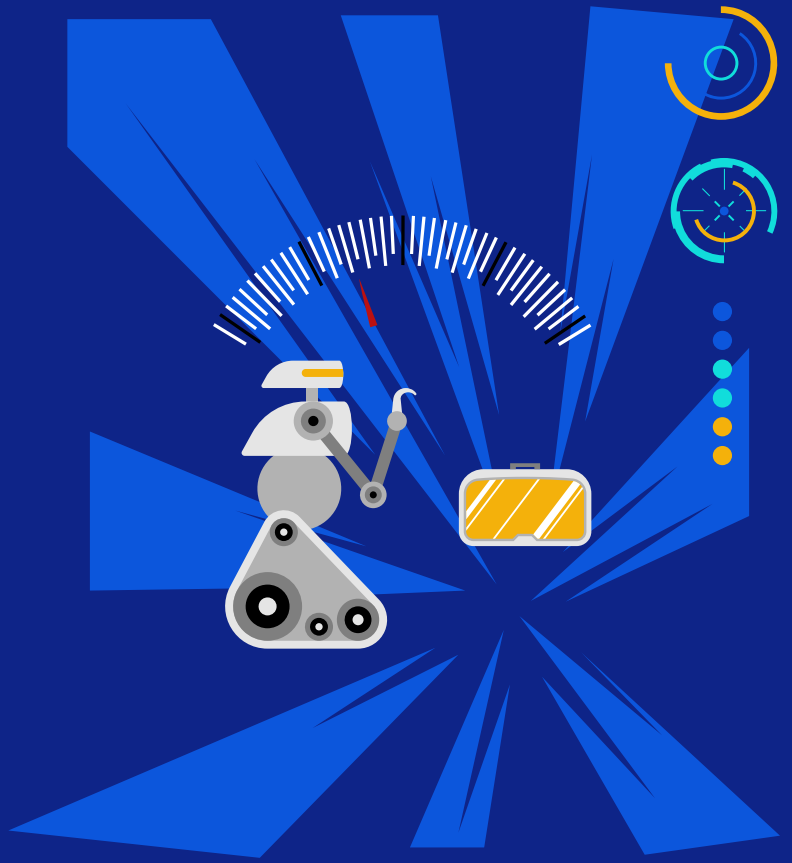

Dataframe

```
1 df = pd.read_csv("adult.data", header = None, names = columnas)
2 df.head()
```

	age	work_class	fnlwgt	education	education_num	marital_status	occupation	relationship	race	sex	capital_gain	capital_loss	hours_per_week	native_country	target
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

Resumen variables numéricas

```
1 df.describe().T
```

	count	mean	std	min	25%	50%	75%	max
age	32561.0	38.581647	13.640433	17.0	28.0	37.0	48.0	90.0
fnlwgt	32561.0	189778.366512	105549.977697	12285.0	117827.0	178356.0	237051.0	1484705.0
education_num	32561.0	10.080679	2.572720	1.0	9.0	10.0	12.0	16.0
capital_gain	32561.0	1077.648844	7385.292085	0.0	0.0	0.0	0.0	99999.0
capital_loss	32561.0	87.303830	402.960219	0.0	0.0	0.0	0.0	4356.0
hours_per_week	32561.0	40.437456	12.347429	1.0	40.0	40.0	45.0	99.0

Scatter Matrix

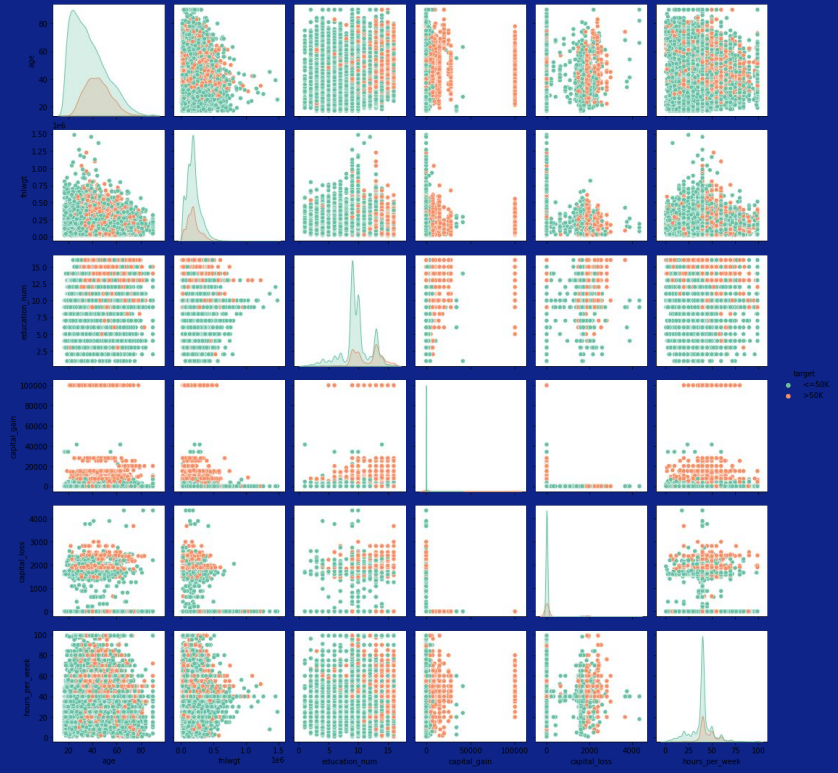

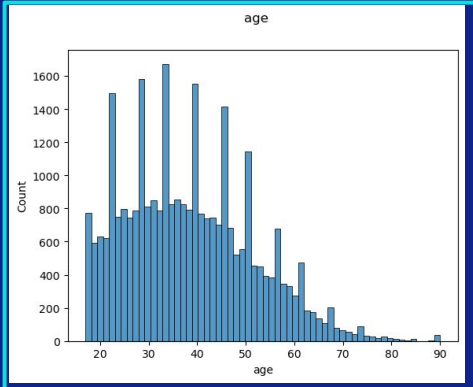
age
Count
age

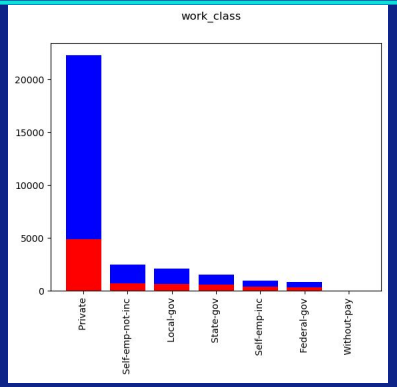
work_class
Private
Self-emp-not-inc
Local-gov
State-gov
Self-emp-inc
Federal-gov
Without-pay

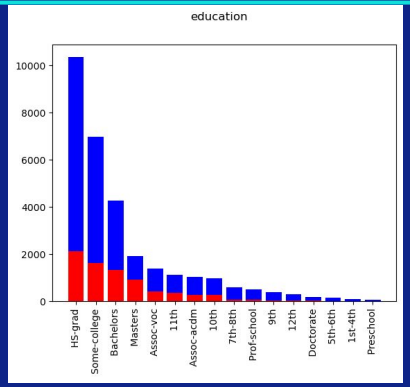
education
HS-grad
Some-college
Bachelors
Masters
Assoc-voc
11th
Assoc-acdm
10th
7th-8th
Prof-school
9th
12th
Doctorate
5th-6th
1st-4th
Preschool

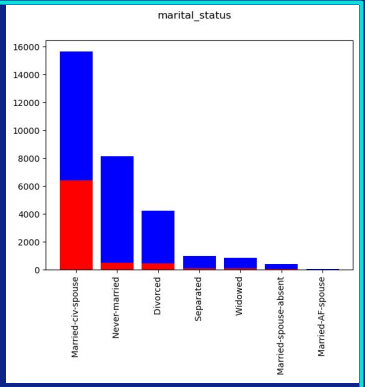
marital_status
Married-civ-spouse
Never-married
Divorced
Separated
Widowed
Married-spouse-absent
Married-AF-spouse

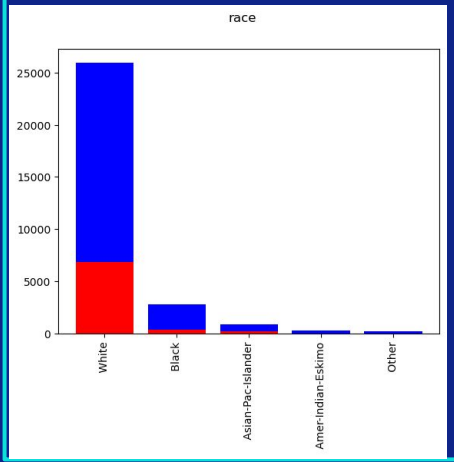
race
White
Black
Asian-Pac-Islander
Amer-Indian-Eskimo
Other

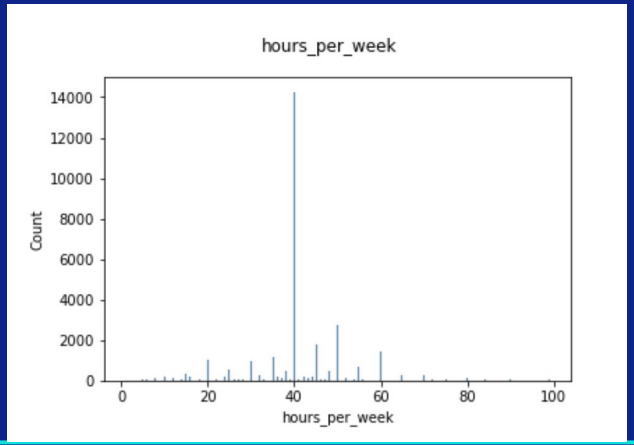
hours_per_week
Count
hours_per_week

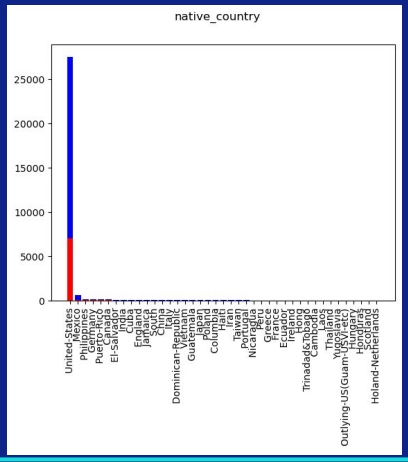
native_country
United-States
Mexico
Philippines
Germany
Puerto-Rico
Canada
El-Salvador
India
Cuba
England
Jamaica
South
China
Italy
Dominican-Republic
Vietnam
Guatemala
Japan
Poland
Columbia
Iran
Taiwan
Haiti
Portugal
Nicaragua
Peru
France
Greece
Ecuador
Ireland
Hong
Trinadad&Tobago
Cambodia
Laos
Thailand
Yugoslavia
Outlying-US(Guam-USVI-etc)
Hungary
Honduras
Scotland
Holand-Netherlands

03

Preprocesado

Eliminamos

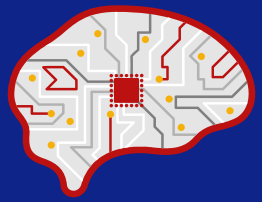

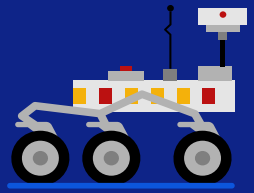

capital_gain

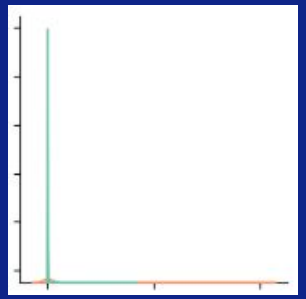

capital_loss

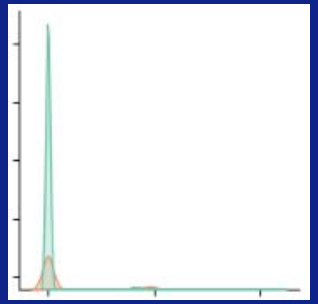

education_num

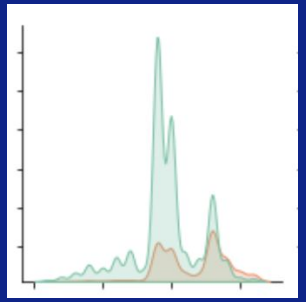

fnlwgt

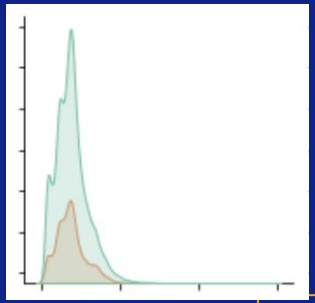

Agrupamos

work_class

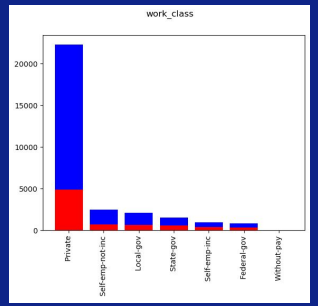

"Private", "Other"

race

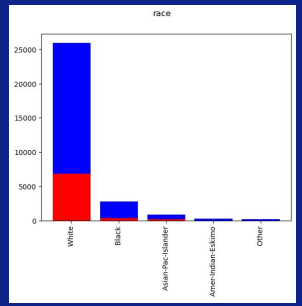

"White", "Other"

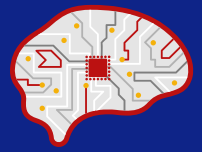

native_country

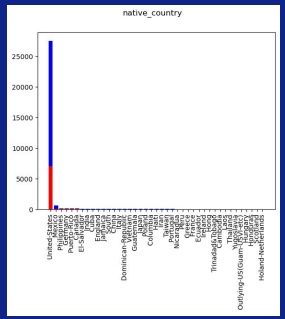

"United-States", "Other"

OneHotEncodeamos

education

marital_status

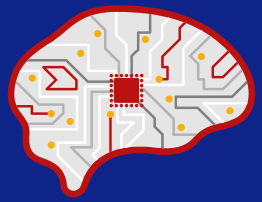

occupation

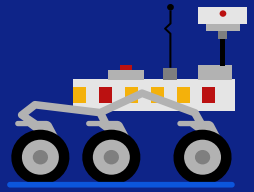

relationship

```
jj = OneHotEncoder(drop= 'first')
```

```
encoded_df = pd.DataFrame(jj.fit_transform(df_semiclean[['education',
                          'marital_status', 'occupation', 'relationship']]).toarray())
```

Matriz de correlación

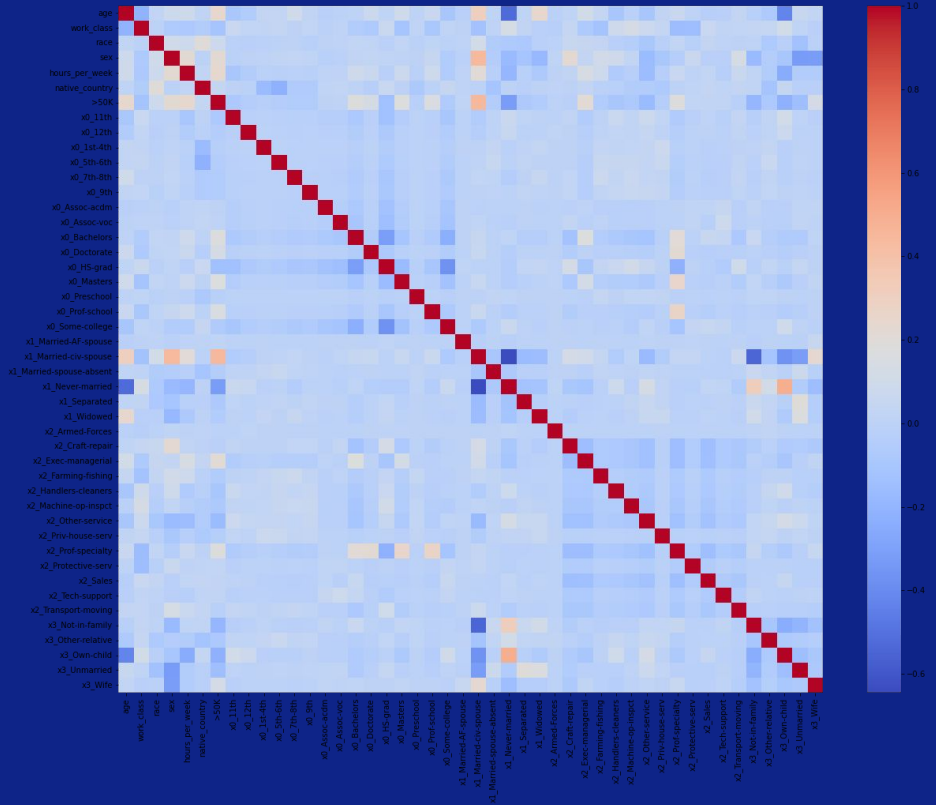

Train y Test split

```
y = df_final['>50K']
X = df_final.drop('>50K', axis= 1)
X_train_sn,  X_test_sn, y_train_sn, y_test_sn = train_test_split(X, y, train_size=0.8, random_state=1729)
```

StandardScaler()

```
X_train_std,  X_test_std, y_train_std, y_test_std = train_test_split(X_std, y_std, train_size=0.8, random_state=1729)

scaler = preprocessing.StandardScaler().fit(X_train_std[columns_name])
X_train_std[columns_name] = scaler.transform(X_train_std[columns_name])

scaler = preprocessing.StandardScaler().fit(X_test_std[columns_name])
X_test_std[columns_name] = scaler.transform(X_test_std[columns_name])
```

Dos datasets

Sin estandarizar

Estandarizado

04

Modelos

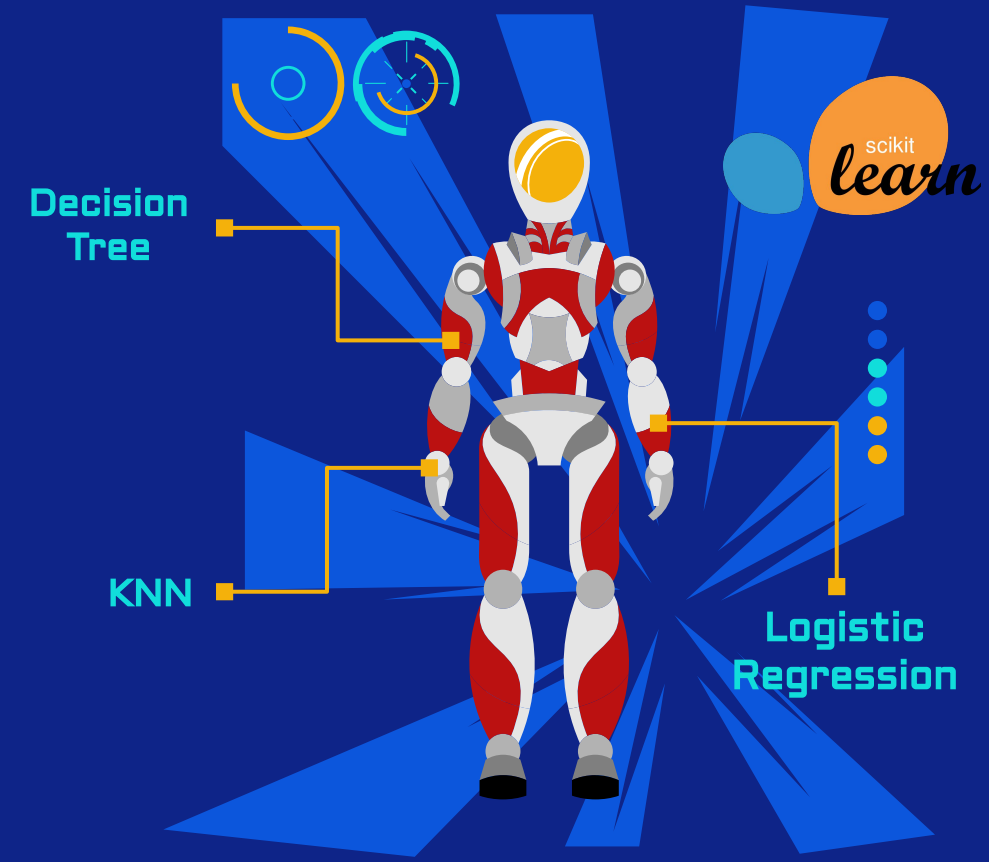

Regresión Logística

Regresión logística SN train

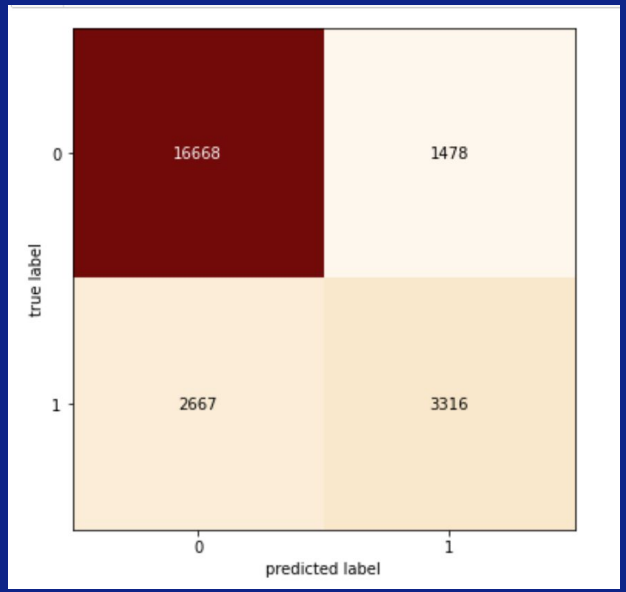

Regresión logística SN test

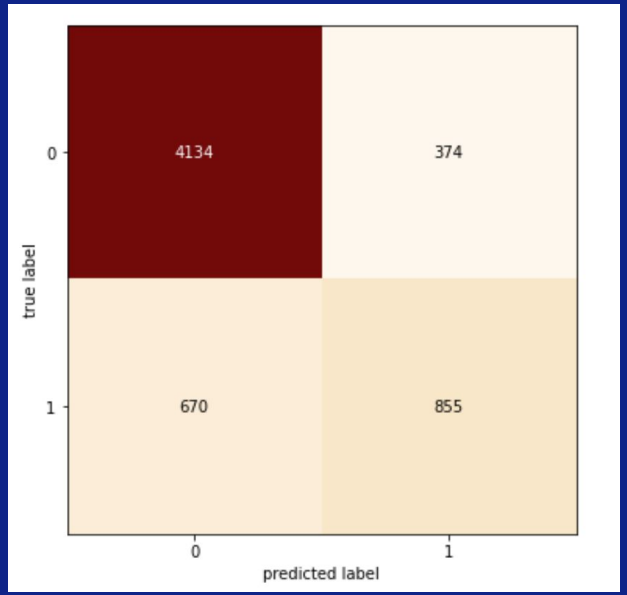

Regresión Logística

Regresión logística STD train

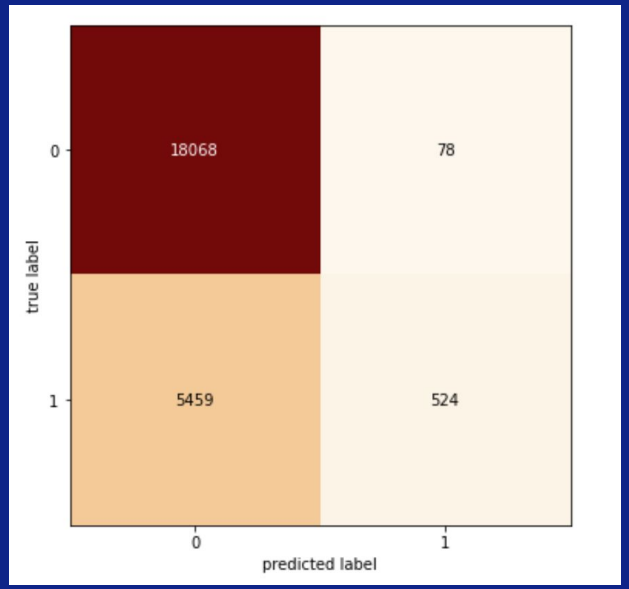

Regresión logística STD test

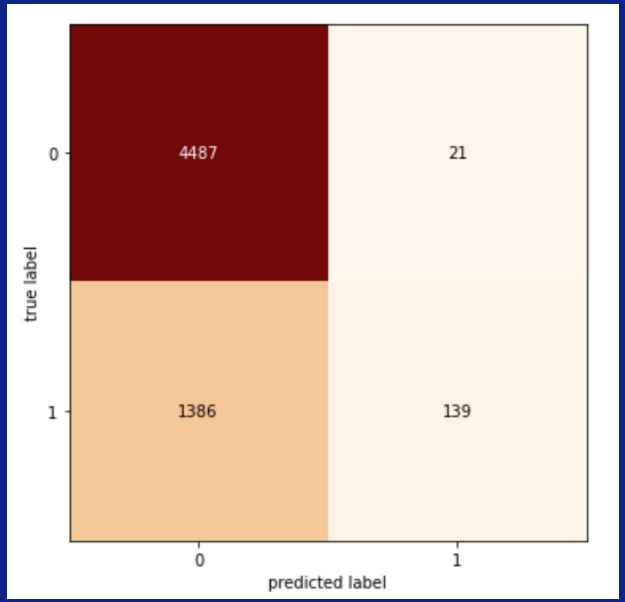

KNN not normalized

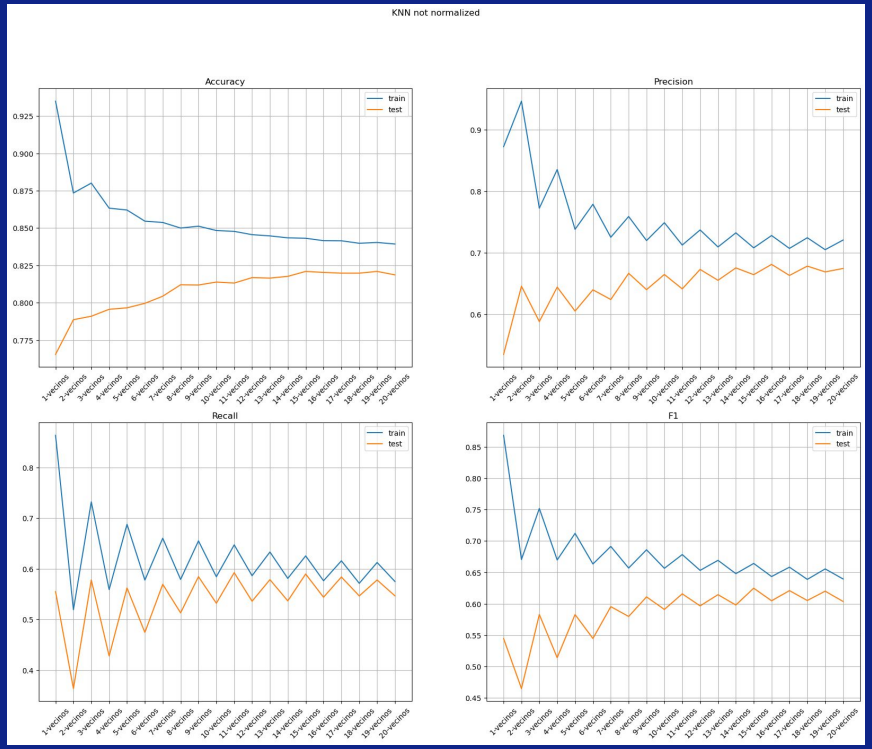

KNN Normalized

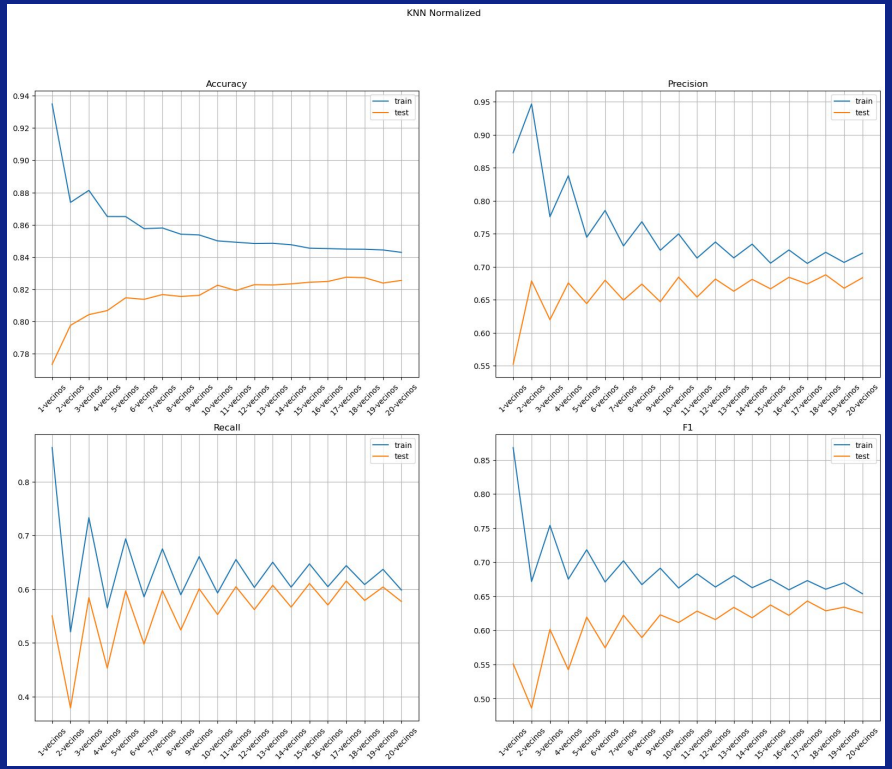

Decision Tree not normalized

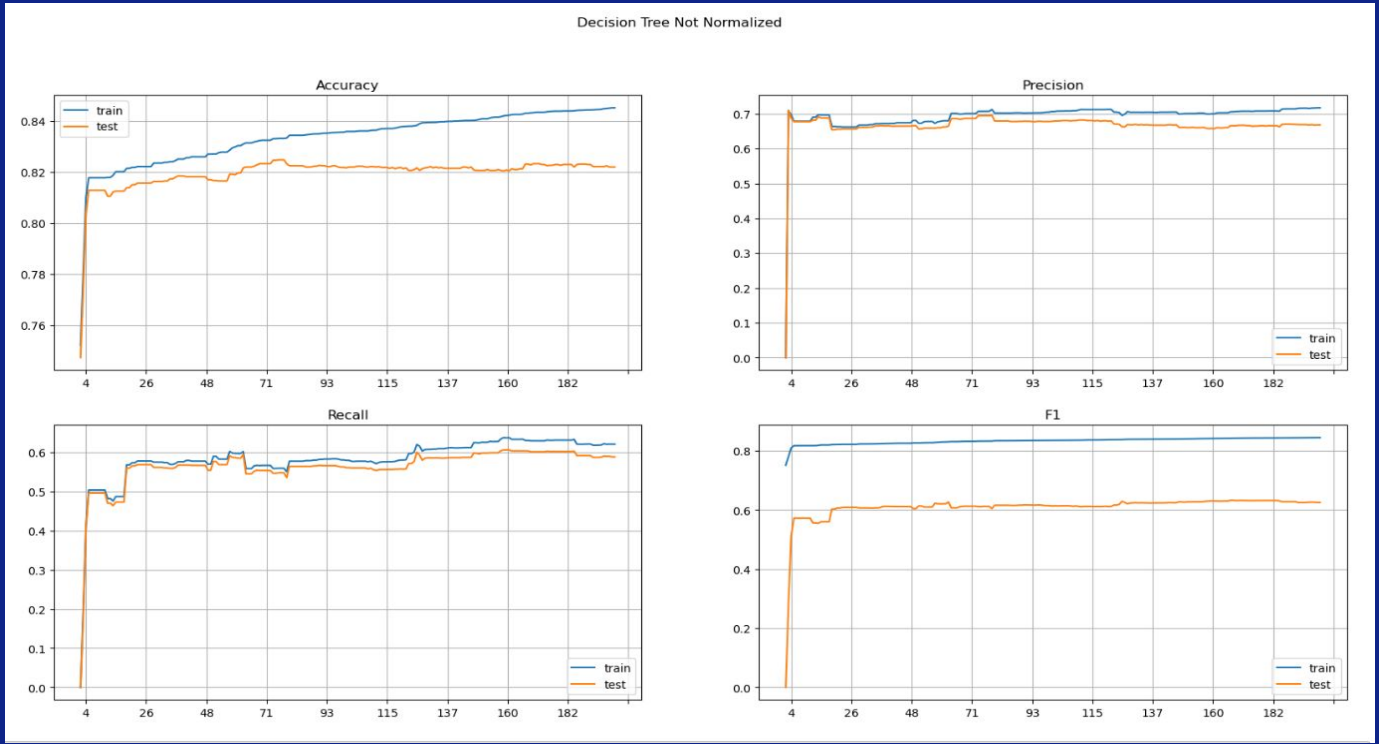

Decision Tree normalized

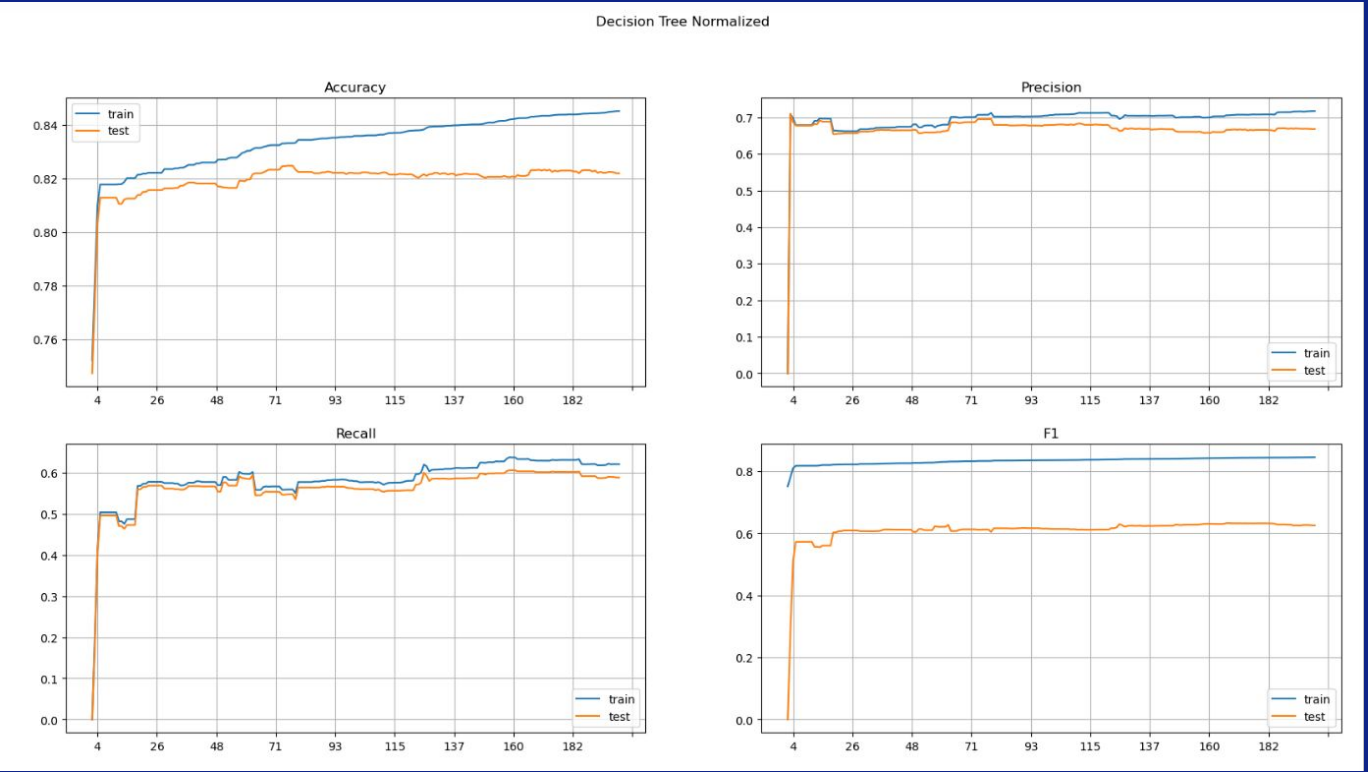

Decision Tree

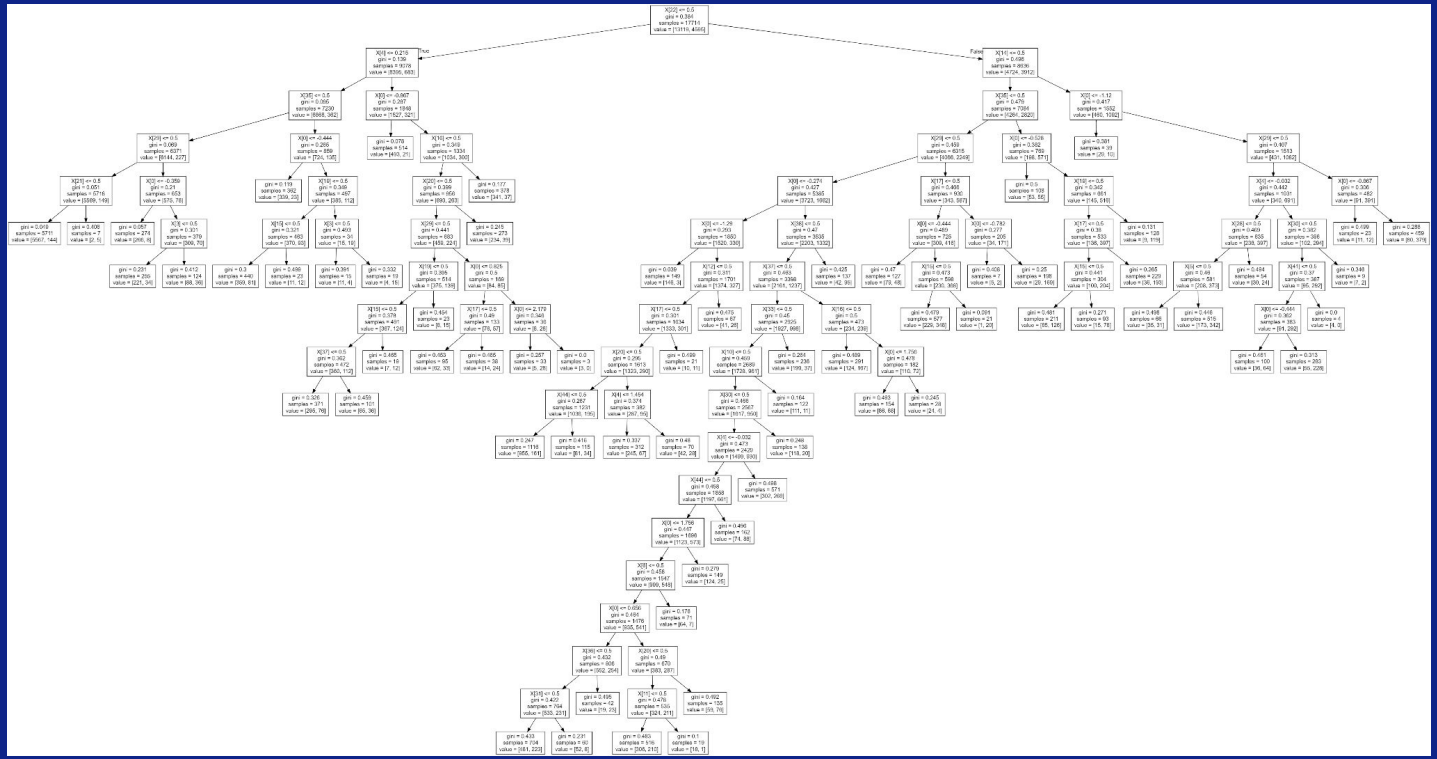

Curvas ROC

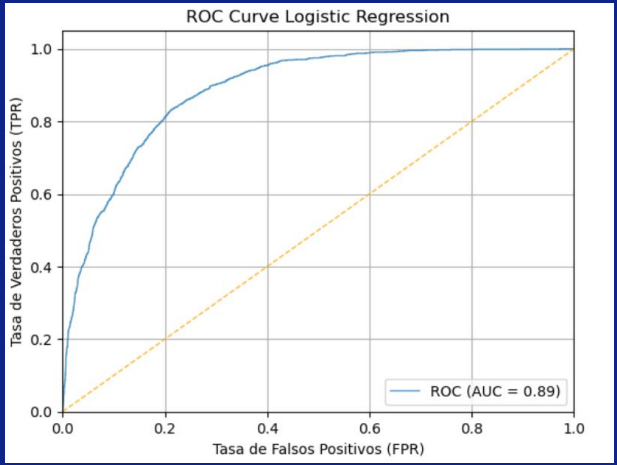
ROC Curve Logistic Regression
Tasa de Verdaderos Positivos (TPR)
Tasa de Falsos Positivos (FPR)
ROC (AUC = 0.89)
0.0
0.2
0.4
0.6
0.8
1.0

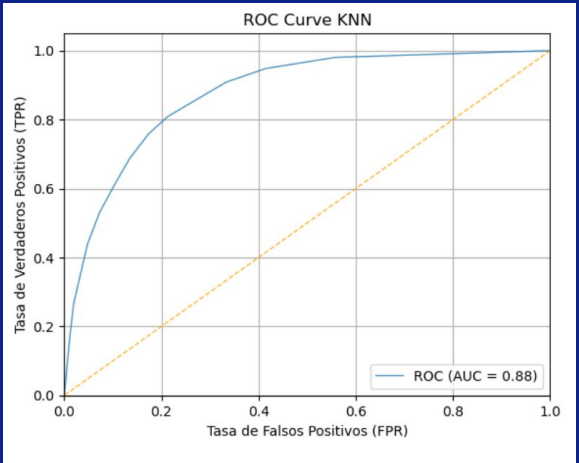
ROC Curve KNN
Tasa de Verdaderos Positivos (TPR)
Tasa de Falsos Positivos (FPR)
ROC (AUC = 0.88)
0.0
0.2
0.4
0.6
0.8
1.0

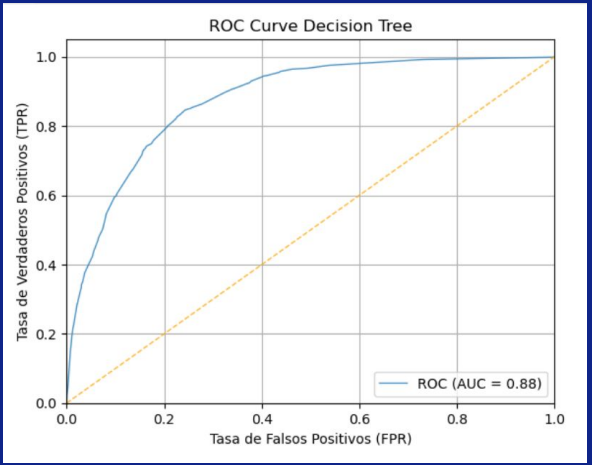
ROC Curve Decision Tree
Tasa de Verdaderos Positivos (TPR)
Tasa de Falsos Positivos (FPR)
ROC (AUC = 0.88)
0.0
0.2
0.4
0.6
0.8
1.0

Comparación de Resultados

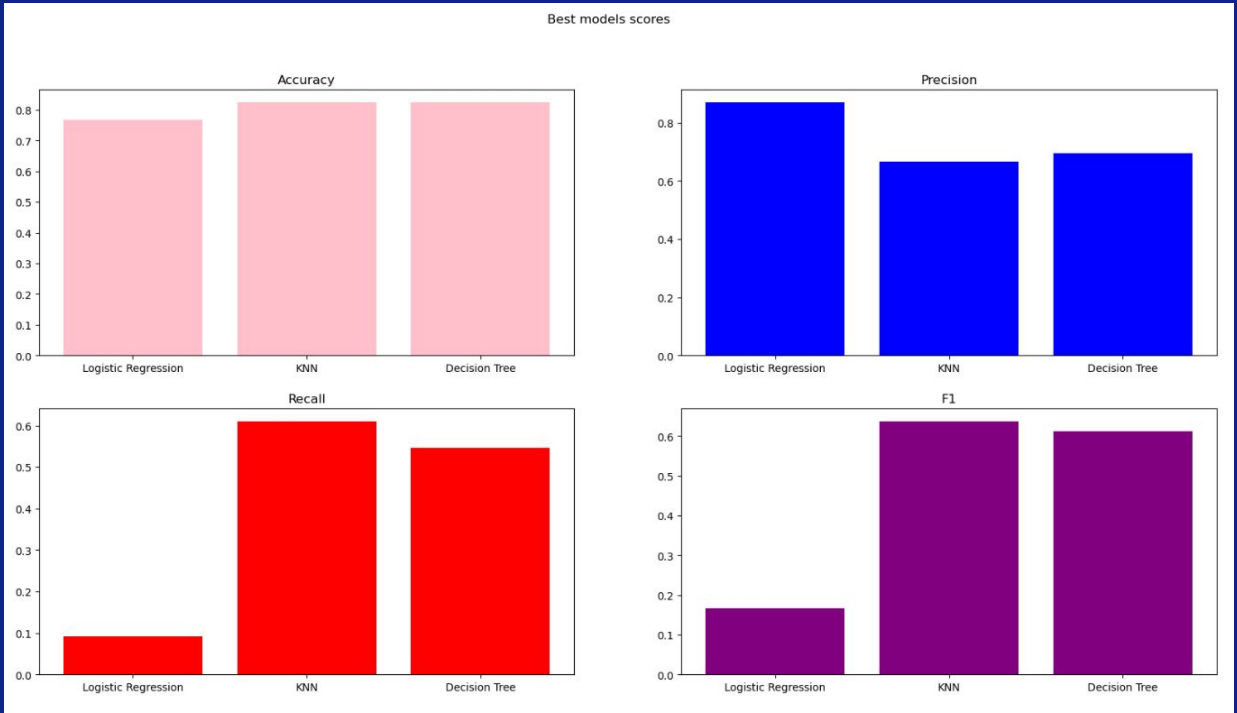

Preguntas